

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
from keras.datasets import fashion_mnist, cifar10
from keras.models import Sequential
from keras.layers import Dense, Activation, Dropout, Flatten, Conv2D, MaxPooling2D
from keras.layers.normalization import BatchNormalization
warnings.filterwarnings("ignore")
```

```
In [2]: (xtrain, ytrain), (xtest, ytest) = cifar10.load_data()
class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>
(<https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>)
170500096/170498071 [=====] - 2s 0us/step

```
In [3]: print("Total Number of images :- ", xtrain.shape[0] + ytrain.shape[0])
```

Total Number of images :- 100000

```
In [4]: fig = plt.figure()
_, axs = plt.subplots(10,10, figsize=(25,25))
axs = axs.flatten()
for img, ax,k in zip(xtrain, axs,ytrain):
    ax.axis("off")
    ax.set_title(class_names[k[0]],fontsize=15)
    ax.imshow(img)
plt.suptitle('X - Train Data',fontsize=25)
plt.show()
```

```
In [5]: from sklearn.model_selection import train_test_split
xtrain,xval,ytrain,yval=train_test_split(xtrain,ytrain,test_size=.3)

#Dimension of the CIFAR10 dataset
print((xtrain.shape,ytrain.shape))
print((xval.shape,yval.shape))
print((xtest.shape,ytest.shape))

((35000, 32, 32, 3), (35000, 1))
((15000, 32, 32, 3), (15000, 1))
((10000, 32, 32, 3), (10000, 1))
```

```
In [6]: from sklearn.utils.multiclass import unique_labels
from keras.utils import to_categorical


#Since we have 10 classes we should expect the shape[1] of y_train,y_val and y_test
ytrain=to_categorical(ytrain)
yval=to_categorical(yval)
ytest=to_categorical(ytest)

#Verifying the dimension after one hot encoding
print((xtrain.shape,ytrain.shape))
print((xval.shape,yval.shape))
print((xtest.shape,ytest.shape))

((35000, 32, 32, 3), (35000, 10))
((15000, 32, 32, 3), (15000, 10))
((10000, 32, 32, 3), (10000, 10))
```

```
In [7]: from keras.preprocessing.image import ImageDataGenerator
train_generator = ImageDataGenerator(rotation_range=2, horizontal_flip=True,zoom_range=0.1)
val_generator = ImageDataGenerator(rotation_range=2, horizontal_flip=True,zoom_range=0.1)
test_generator = ImageDataGenerator(rotation_range=2, horizontal_flip=True,zoom_range=0.1)
#Fitting the augmentation defined above to the data
train_generator.fit(xtrain)
val_generator.fit(xval)
test_generator.fit(xtest)
```

```
In [8]: ann = Sequential()
x = Conv2D(filters=64,kernel_size=(5,5),input_shape=(32,32,3))
ann.add(x)
x1w = x.get_weights()[0][:,:,0,:]
for i in range(1,26):
    plt.subplot(5,5,i)
    plt.imshow(x1w[:,:,:],interpolation="nearest",cmap="gray")
plt.show()
```



```
In [9]: AlexNet = Sequential()

#1st Convolutional Layer
AlexNet.add(Conv2D(filters=96, input_shape=(32,32,3), kernel_size=(11,11), stride
AlexNet.add(BatchNormalization())
AlexNet.add(Activation('relu'))
AlexNet.add(MaxPooling2D(pool_size=(2,2), strides=(2,2), padding='same'))

#2nd Convolutional Layer
AlexNet.add(Conv2D(filters=256, kernel_size=(5, 5), strides=(1,1), padding='same')
AlexNet.add(BatchNormalization())
AlexNet.add(Activation('relu'))
AlexNet.add(MaxPooling2D(pool_size=(2,2), strides=(2,2), padding='same'))

#3rd Convolutional Layer
AlexNet.add(Conv2D(filters=384, kernel_size=(3,3), strides=(1,1), padding='same')
AlexNet.add(BatchNormalization())
AlexNet.add(Activation('relu'))

#4th Convolutional Layer
AlexNet.add(Conv2D(filters=384, kernel_size=(3,3), strides=(1,1), padding='same')
AlexNet.add(BatchNormalization())
AlexNet.add(Activation('relu'))

#5th Convolutional Layer
AlexNet.add(Conv2D(filters=256, kernel_size=(3,3), strides=(1,1), padding='same')
AlexNet.add(BatchNormalization())
AlexNet.add(Activation('relu'))
AlexNet.add(MaxPooling2D(pool_size=(2,2), strides=(2,2), padding='same'))

#Passing it to a Fully Connected Layer
AlexNet.add(Flatten())
# 1st Fully Connected Layer
AlexNet.add(Dense(4096, input_shape=(32,32,3)))
AlexNet.add(BatchNormalization())
AlexNet.add(Activation('relu'))
# Add Dropout to prevent overfitting
AlexNet.add(Dropout(0.4))

#2nd Fully Connected Layer
AlexNet.add(Dense(4096))
AlexNet.add(BatchNormalization())
AlexNet.add(Activation('relu'))
#Add Dropout
AlexNet.add(Dropout(0.4))

#3rd Fully Connected Layer
AlexNet.add(Dense(1000))
AlexNet.add(BatchNormalization())
AlexNet.add(Activation('relu'))
#Add Dropout
AlexNet.add(Dropout(0.4))

#Output Layer
AlexNet.add(Dense(10))
AlexNet.add(BatchNormalization())
```

```
AlexNet.add(Activation('softmax'))
```

```
#Model Summary
```

```
AlexNet.summary()
```

```
Model: "sequential_1"
```

Layer (type)	Output Shape	Param #
=====		
conv2d_1 (Conv2D)	(None, 8, 8, 96)	34944
batch_normalization (Batch Normalization)	(None, 8, 8, 96)	384
activation (Activation)	(None, 8, 8, 96)	0
max_pooling2d (MaxPooling2D)	(None, 4, 4, 96)	0
conv2d_2 (Conv2D)	(None, 4, 4, 256)	614656
batch_normalization_1 (Batch Normalization)	(None, 4, 4, 256)	1024
activation_1 (Activation)	(None, 4, 4, 256)	0
max_pooling2d_1 (MaxPooling2D)	(None, 2, 2, 256)	0
conv2d_3 (Conv2D)	(None, 2, 2, 384)	885120
batch_normalization_2 (Batch Normalization)	(None, 2, 2, 384)	1536
activation_2 (Activation)	(None, 2, 2, 384)	0
conv2d_4 (Conv2D)	(None, 2, 2, 384)	1327488
batch_normalization_3 (Batch Normalization)	(None, 2, 2, 384)	1536
activation_3 (Activation)	(None, 2, 2, 384)	0
conv2d_5 (Conv2D)	(None, 2, 2, 256)	884992
batch_normalization_4 (Batch Normalization)	(None, 2, 2, 256)	1024
activation_4 (Activation)	(None, 2, 2, 256)	0
max_pooling2d_2 (MaxPooling2D)	(None, 1, 1, 256)	0
flatten (Flatten)	(None, 256)	0
dense (Dense)	(None, 4096)	1052672
batch_normalization_5 (Batch Normalization)	(None, 4096)	16384
activation_5 (Activation)	(None, 4096)	0
dropout (Dropout)	(None, 4096)	0
dense_1 (Dense)	(None, 4096)	16781312
batch_normalization_6 (Batch Normalization)	(None, 4096)	16384

activation_6 (Activation)	(None, 4096)	0
dropout_1 (Dropout)	(None, 4096)	0
dense_2 (Dense)	(None, 1000)	4097000
batch_normalization_7 (Batch Normalization)	(None, 1000)	4000
activation_7 (Activation)	(None, 1000)	0
dropout_2 (Dropout)	(None, 1000)	0
dense_3 (Dense)	(None, 10)	10010
batch_normalization_8 (Batch Normalization)	(None, 10)	40
activation_8 (Activation)	(None, 10)	0
=====		
Total params: 25,730,506		
Trainable params: 25,709,350		
Non-trainable params: 21,156		

```
In [10]: AlexNet.compile(loss = 'categorical_crossentropy', optimizer= 'adam', metrics=['a
```

```
In [11]: batch_size= 128
epochs=200
learn_rate=.001
from keras.callbacks import ReduceLROnPlateau
lrr= ReduceLROnPlateau( monitor='val_acc', factor=.01, patience=3, min_lr=
```



```
In [12]: AlexNet.fit_generator(train_generator.flow(xtrain, ytrain, batch_size=batch_size),
                             epochs = epochs, steps_per_epoch = 5,
                             validation_data = val_generator.flow(xval, yval, batch_size=batch_size),
                             validation_steps = 5, callbacks = [lrr], verbose=1)
```

Epoch 1/200

5/5 [=====] - 35s 280ms/step - loss: 2.5649 - accuracy: 0.1277 - val_loss: 60.0688 - val_accuracy: 0.0969

WARNING:tensorflow:Learning rate reduction is conditioned on metric `val_acc` which is not available. Available metrics are: loss,accuracy,val_loss,val_accuracy,lr

Epoch 2/200

5/5 [=====] - 1s 144ms/step - loss: 2.1025 - accuracy: 0.2146 - val_loss: 111.7256 - val_accuracy: 0.0797

WARNING:tensorflow:Learning rate reduction is conditioned on metric `val_acc` which is not available. Available metrics are: loss,accuracy,val_loss,val_accuracy,lr

Epoch 3/200

5/5 [=====] - 1s 147ms/step - loss: 1.9375 - accuracy: 0.2368 - val_loss: 119.3911 - val_accuracy: 0.0906

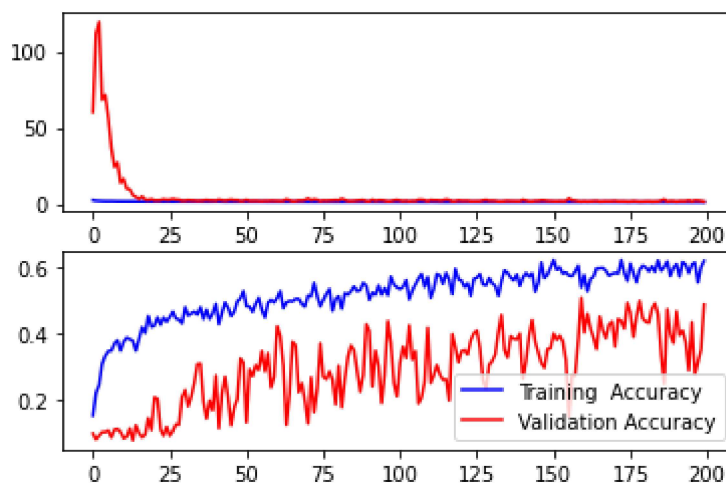
WARNING:tensorflow:Learning rate reduction is conditioned on metric `val_acc` which is not available. Available metrics are: loss,accuracy,val_loss,val_accuracy,lr

Epoch 4/200

```
In [13]: # AlexNet.history.history.L
```

```
In [14]: import matplotlib.pyplot as plt
#Plotting the training and validation loss
f,ax=plt.subplots(2,1) #Creates 2 subplots under 1 column
#Assigning the first subplot to graph training loss and validation loss
ax[0].plot(AlexNet.history.history['loss'],color='b',label='Training Loss')
ax[0].plot(AlexNet.history.history['val_loss'],color='r',label='Validation Loss')
#Plotting the training accuracy and validation accuracy
ax[1].plot(AlexNet.history.history['accuracy'],color='b',label='Training Accuracy')
ax[1].plot(AlexNet.history.history['val_accuracy'],color='r',label='Validation Accuracy')
plt.legend()
```

```
Out[14]: <matplotlib.legend.Legend at 0x7f9300564dd0>
```




```
In [15]: predictions = AlexNet.predict_classes(xtest)
ytest = np.argmax(ytest,axis=1)
```

```
In [16]: from sklearn.metrics import confusion_matrix,plot_confusion_matrix,accuracy_score
print(confusion_matrix(predictions,ytest))
```

```
[[ 776  13 302  90 176  48  30  63 163  31]
 [129 937  94 109  71  60 116  38 210 587]
 [  1  0  66  1 12  3  6  0  0  0]
 [  5  0  39  75 52 20 59  7  2  5]
 [  0  0  9  1 68  0  3  0  0  0]
 [ 25 11 357 576 313 771 191 228 27 17]
 [  2  2  45  51 44 18 538  2  0  6]
 [  6  3  49  32 210 45 27 608  2 11]
 [ 31  9 18 30 28 20 10  5 567 19]
 [ 25 25 21 35 26 15 20 49 29 324]]
```

```
In [17]: print(accuracy_score(predictions,ytest))
```

```
0.473
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

In []:

In []:

In []:

In []:

In []:

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from keras.layers import AveragePooling2D, Conv2D, Dense, Activation, Flatten
from keras.models import Sequential
from keras.utils import to_categorical
from sklearn.metrics import accuracy_score, confusion_matrix
from keras.datasets import cifar10
```

```
In [2]: (xtrain,ytrain),(xtest,ytest) = cifar10.load_data()
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>
(<https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>)
170500096/170498071 [=====] - 4s 0us/step

```
In [3]: ytrain =to_categorical(ytrain)
ytest = to_categorical(ytest)
```

```
In [4]: model = Sequential()
model.add(Conv2D(64,kernel_size=(3,3),strides = (2,2),input_shape=xtrain.shape[1:]))
model.add(AveragePooling2D((2,2)))
model.add(Conv2D(20,(3,3)))
model.add(AveragePooling2D((2,2)))
model.add(Flatten())
model.add(Dense(16))
model.add(Dense(12))
model.add(Dense(10,))
model.add(Activation('softmax'))
model.compile(loss = 'categorical_crossentropy', optimizer= 'adam', metrics=['acc'])
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 15, 15, 64)	1792
=====		
average_pooling2d (AveragePo	(None, 7, 7, 64)	0
=====		
conv2d_1 (Conv2D)	(None, 5, 5, 20)	11540
=====		
average_pooling2d_1 (Average	(None, 2, 2, 20)	0
=====		
flatten (Flatten)	(None, 80)	0
=====		
dense (Dense)	(None, 16)	1296
=====		
dense_1 (Dense)	(None, 12)	204
=====		
dense_2 (Dense)	(None, 10)	130
=====		
activation (Activation)	(None, 10)	0
=====		
Total params: 14,962		
Trainable params: 14,962		
Non-trainable params: 0		
=====		

In [5]: `history = model.fit(xtrain,ytrain,batch_size=128,epochs=200)`

```
Epoch 1/200
391/391 [=====] - 34s 4ms/step - loss: 11.5363 - acc
uracy: 0.1779
Epoch 2/200
391/391 [=====] - 2s 4ms/step - loss: 2.1525 - accur
acy: 0.2728
Epoch 3/200
391/391 [=====] - 2s 4ms/step - loss: 1.9921 - accur
acy: 0.3067
Epoch 4/200
391/391 [=====] - 2s 4ms/step - loss: 1.9247 - accur
acy: 0.3289
Epoch 5/200
391/391 [=====] - 2s 4ms/step - loss: 1.8791 - accur
acy: 0.3426
Epoch 6/200
391/391 [=====] - 2s 4ms/step - loss: 1.8528 - accur
acy: 0.3568
Epoch 7/200
391/391 [=====] - 2s 4ms/step - loss: 1.8370 - accur
acy: 0.3670
```


In [6]: `history = pd.DataFrame(history.history)`
`history`

Out[6]:

	loss	accuracy
0	4.973685	0.21240
1	2.101855	0.28438
2	1.959686	0.31452
3	1.909298	0.33428
4	1.876671	0.34332
...
195	1.764839	0.39284
196	1.764209	0.39386
197	1.764538	0.39380
198	1.763331	0.39296
199	1.762880	0.39166

200 rows × 2 columns

```
In [7]: plt.figure(figsize=(16,4))
plt.plot(history["accuracy"],color="red",label="accuracy")
plt.plot(history["loss"],color="blue",label="loss")
plt.legend()
plt.grid()
plt.show()
```



```
In [8]: predictions = np.argmax(model.predict(xtest),axis=1)
```

```
In [9]: ytest = np.argmax(ytest,axis=1)
```

```
In [10]: print(accuracy_score(predictions,ytest))
```

0.3744

```
In [11]: print(confusion_matrix(predictions,ytest))
```

```
[[377  58  69  31  37  37  20  39 127  54]
 [ 26 381  32  34  13  33  19  32  38 125]
 [ 71  41 284 137 146 125  99  83  22  26]
 [ 22  46  67 265  71 165 133  70  31  22]
 [ 19  41 182  89 365 120 152 114  19  28]
 [ 19  46  67 128  61 279  55  57  42  30]
 [ 23  25  92 115 116  72 410  37   9  42]
 [ 47  41  83  52  95  58  48 391  26  47]
 [296 128  81  71  48  71  28  60 515 149]
 [100 193  43  78  48  40  36 117 171 477]]
```

```
In [11]:
```