

Name: Prekshita vasudeo patil

registration No.: 20MAI0073

Assignment-4

Task-2 ¶

In [1]:

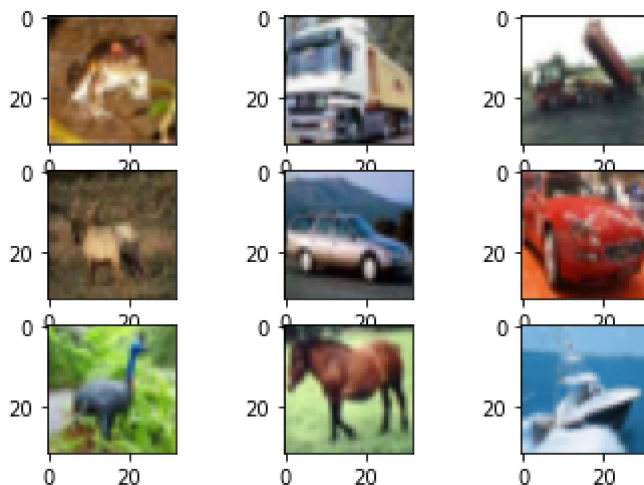
```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import seaborn as sns
import keras as k
import tensorflow as tf
from keras.datasets import cifar10
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from sklearn.model_selection import train_test_split
```

In [2]:

```
(trainX, trainy), (testX, testy) = cifar10.load_data()
print('Train: X=%s, y=%s' % (trainX.shape, trainy.shape))
print('Test: X=%s, y=%s' % (testX.shape, testy.shape))
for i in range(9):
    plt.subplot(330 + 1 + i)
    plt.imshow(trainX[i])
plt.show()
```

Train: X=(50000, 32, 32, 3), y=(50000, 1)

Test: X=(10000, 32, 32, 3), y=(10000, 1)



```
In [3]: from keras.utils import to_categorical
trainy = to_categorical(trainy)
testy = to_categorical(testy)
```

```
In [4]: from keras.layers import Conv2D,MaxPooling2D,Flatten,Dense,BatchNormalization
from keras.models import Sequential
from keras.optimizers import SGD
from keras.preprocessing.image import ImageDataGenerator
```

```
In [5]: datagen = ImageDataGenerator(width_shift_range=0.1, height_shift_range=0.1, hori
```

```
In [6]: model = Sequential()
model.add(Conv2D(32, (3, 3), activation='relu', kernel_initializer='he_uniform',
model.add(Conv2D(32, (3, 3), activation='relu', kernel_initializer='he_uniform',
model.add(MaxPooling2D((2, 2)))
model.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform',
model.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform',
model.add(MaxPooling2D((2, 2)))
model.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform',
model.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform',
model.add(MaxPooling2D((2, 2)))
model.add(BatchNormalization()) # Adding Batch Normalization
model.add(Flatten())
model.add(Dense(128, activation='relu', kernel_initializer='he_uniform'))
model.add(Dense(10, activation='softmax'))
# compile model
opt = SGD(lr=0.001, momentum=0.9)
model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
conv2d_1 (Conv2D)	(None, 32, 32, 32)	9248
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_2 (Conv2D)	(None, 16, 16, 64)	18496
conv2d_3 (Conv2D)	(None, 16, 16, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 64)	0
conv2d_4 (Conv2D)	(None, 8, 8, 128)	73856
conv2d_5 (Conv2D)	(None, 8, 8, 128)	147584
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 128)	0
batch_normalization (Batch Normalization)	(None, 4, 4, 128)	512
flatten (Flatten)	(None, 2048)	0
dense (Dense)	(None, 128)	262272
dense_1 (Dense)	(None, 10)	1290
Total params: 551,082		
Trainable params: 550,826		
Non-trainable params: 256		

```
In [7]: it_train = datagen.flow(trainX, trainy, batch_size=64)
```


```
In [8]: steps = int(trainX.shape[0] / 64)
        history = model.fit_generator(it_train, steps_per_epoch=steps, epochs=5, val_i

/usr/local/lib/python3.7/dist-packages/tensorflow/python/keras/engine/training.
py:1844: UserWarning: `Model.fit_generator` is deprecated and will be removed i
n a future version. Please use `Model.fit`, which supports generators.
      warnings.warn("`Model.fit_generator` is deprecated and ")
```

```
In [9]: history = pd.DataFrame(history.history)
```

```
In [10]: history.plot.line(figsize=(16,10))
```

```
Out[10]: <matplotlib.axes._subplots.AxesSubplot at 0x7ff56ba54410>
```



```
In [16]: ypred = np.argmax(model.predict(testX),axis=1)
```

```
In [17]: ypred
```

```
Out[17]: array([3, 1, 8, ..., 5, 1, 7])
```

```
In [18]: testty = np.argmax(testty,axis=1)
```

```
In [19]: testty
```

```
Out[19]: array([3, 8, 8, ..., 5, 1, 7])
```

In [20]: `accuracy_score(ypred, testty)`

Out[20]: 0.6449

In []:

In []:

In []:

In []:

In []:

```
# Name: Prekshita vasudeo patil
# registration No.: 20MAI0073
# Assignment-4
# Task-1
# Link: https://github.com/prekshita19/DL-Assignments/tree/main/Assignment-4
```

In []:

```
import pandas as pd
import numpy as np
import tensorflow as tf
import keras
import matplotlib.pyplot as plt
import seaborn as sns
import scipy
from sklearn.model_selection import train_test_split
```

In [2]:

```
(xtrain,ytrain),(xtest,ytest) = keras.datasets.cifar10.load_data()
classes = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse']
```

In [3]:

```
plt.imshow(xtrain[0])
plt.title(classes[ytrain[0][0]],)
plt.axis(False)
plt.show()
```

frog



```
In [4]: # selecting 50% less data from xtrain
(xtrain_new_50,xtest_new_50,ytrain_new_50,ytest_new_50) =train_test_split(xtrain,
xtrain_1,xtest_1,ytrain_1,ytest_1 = train_test_split(xtrain_new_50,ytrain_new_50,
print("Xtrain original :- ",xtrain.shape)
print("Xtrain 50% selected from xtrain:-",xtrain_new_50.shape)
print("70% selected from that 50% training :- ",xtrain_1.shape)
```

```
Xtrain original :- (50000, 32, 32, 3)
Xtrain 50% selected from xtrain:- (25000, 32, 32, 3)
70% selected from that 50% training :- (17500, 32, 32, 3)
```

```
In [5]: from keras.models import Sequential
from keras.layers import Conv2D,Activation,BatchNormalization,MaxPooling2D,Dense,
```

```
In [6]: from keras.utils import to_categorical

ytrain = to_categorical(ytrain,10)
ytrain_new_50 = to_categorical(ytrain_new_50,10)
ytrain_1 = to_categorical(ytrain_1,10)

ytest = to_categorical(ytest,10)
ytest_new_50 = to_categorical(ytest_new_50,10)
ytest_1 = to_categorical(ytest_1,10)
```

```
In [7]: xtrain_1[0].shape
```

```
Out[7]: (32, 32, 3)
```

```
In [8]: AlexNet = Sequential()

#1st Convolutional Layer
AlexNet.add(Conv2D(filters=96, input_shape=xtrain_1[0].shape, kernel_size=(11,11))
AlexNet.add(BatchNormalization())
AlexNet.add(Activation('relu'))

# 1st Maxpooling Layer
AlexNet.add(MaxPooling2D(pool_size=(2,2), strides=(2,2), padding='same'))

#2nd Convolutional Layer
AlexNet.add(Conv2D(filters=256, kernel_size=(5, 5), strides=(1,1), padding='same')
AlexNet.add(BatchNormalization())
AlexNet.add(Activation('relu'))

# 2nd Maxpooling Layer
AlexNet.add(MaxPooling2D(pool_size=(2,2), strides=(2,2), padding='same'))

#3rd Convolutional Layer
AlexNet.add(Conv2D(filters=384, kernel_size=(3,3), strides=(1,1), padding='same')
AlexNet.add(BatchNormalization())
AlexNet.add(Activation('relu'))

#4th Convolutional Layer
AlexNet.add(Conv2D(filters=384, kernel_size=(3,3), strides=(1,1), padding='same')
AlexNet.add(BatchNormalization())
AlexNet.add(Activation('relu'))

#5th Convolutional Layer
AlexNet.add(Conv2D(filters=256, kernel_size=(3,3), strides=(1,1), padding='same')
AlexNet.add(BatchNormalization())
AlexNet.add(Activation('relu'))

# 3rd Maxpooling Layer
AlexNet.add(MaxPooling2D(pool_size=(2,2), strides=(2,2), padding='same'))

#Passing it to a Fully Connected Layer
AlexNet.add(Flatten())

# 1st Fully Connected Layer
AlexNet.add(Dense(4096, input_shape=(32,32,3)))
AlexNet.add(BatchNormalization())
AlexNet.add(Activation('relu'))

# Add Dropout to prevent overfitting
AlexNet.add(Dropout(0.4))

#2nd Fully Connected Layer
AlexNet.add(Dense(4096))
AlexNet.add(BatchNormalization())
AlexNet.add(Activation('relu'))
#Add Dropout
AlexNet.add(Dropout(0.4))

#3rd Fully Connected Layer
# AlexNet.add(Dense(1000))
```



```

AlexNet.add(BatchNormalization())
AlexNet.add(Activation('relu'))
#Add Dropout
AlexNet.add(Dropout(0.4))

#Output Layer
AlexNet.add(Dense(10))
AlexNet.add(BatchNormalization())
AlexNet.add(Activation('softmax'))

#Model Summary
AlexNet.summary()
# https://cs231n.github.io/transfer-learning/

```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 8, 8, 96)	34944
batch_normalization (Batch Normalization)	(None, 8, 8, 96)	384
activation (Activation)	(None, 8, 8, 96)	0
max_pooling2d (MaxPooling2D)	(None, 4, 4, 96)	0
conv2d_1 (Conv2D)	(None, 4, 4, 256)	614656
batch_normalization_1 (Batch Normalization)	(None, 4, 4, 256)	1024
activation_1 (Activation)	(None, 4, 4, 256)	0
max_pooling2d_1 (MaxPooling2D)	(None, 2, 2, 256)	0
conv2d_2 (Conv2D)	(None, 2, 2, 384)	885120
batch_normalization_2 (Batch Normalization)	(None, 2, 2, 384)	1536
activation_2 (Activation)	(None, 2, 2, 384)	0
conv2d_3 (Conv2D)	(None, 2, 2, 384)	1327488
batch_normalization_3 (Batch Normalization)	(None, 2, 2, 384)	1536
activation_3 (Activation)	(None, 2, 2, 384)	0
conv2d_4 (Conv2D)	(None, 2, 2, 256)	884992
batch_normalization_4 (Batch Normalization)	(None, 2, 2, 256)	1024
activation_4 (Activation)	(None, 2, 2, 256)	0
max_pooling2d_2 (MaxPooling2D)	(None, 1, 1, 256)	0
flatten (Flatten)	(None, 256)	0
dense (Dense)	(None, 4096)	1052672

batch_normalization_5 (Batch Normalization)	(None, 4096)	16384
activation_5 (Activation)	(None, 4096)	0
dropout (Dropout)	(None, 4096)	0
dense_1 (Dense)	(None, 4096)	16781312
batch_normalization_6 (Batch Normalization)	(None, 4096)	16384
activation_6 (Activation)	(None, 4096)	0
dropout_1 (Dropout)	(None, 4096)	0
batch_normalization_7 (Batch Normalization)	(None, 4096)	16384
activation_7 (Activation)	(None, 4096)	0
dropout_2 (Dropout)	(None, 4096)	0
dense_2 (Dense)	(None, 10)	40970
batch_normalization_8 (Batch Normalization)	(None, 10)	40
activation_8 (Activation)	(None, 10)	0
=====		
Total params: 21,676,850		
Trainable params: 21,649,502		
Non-trainable params: 27,348		

In [9]: AlexNet.compile(loss = 'categorical_crossentropy', optimizer= 'adam', metrics=['a

```
In [10]: # checkpoint = ModelCheckpoint('AlexNet.h5',save_best_only=True, monitor='val_acc
history = AlexNet.fit(xtrain_1,ytrain_1,epochs=10,batch_size=32)
```

```
Epoch 1/10
547/547 [=====] - 11s 13ms/step - loss: 1.9141 - accur
acy: 0.3003
Epoch 2/10
547/547 [=====] - 7s 12ms/step - loss: 1.5539 - accura
cy: 0.4400
Epoch 3/10
547/547 [=====] - 7s 12ms/step - loss: 1.4151 - accura
cy: 0.5054
Epoch 4/10
547/547 [=====] - 7s 12ms/step - loss: 1.3250 - accura
cy: 0.5396
Epoch 5/10
547/547 [=====] - 7s 12ms/step - loss: 1.2156 - accura
cy: 0.5769
Epoch 6/10
547/547 [=====] - 7s 12ms/step - loss: 1.1364 - accura
cy: 0.6088
Epoch 7/10
547/547 [=====] - 7s 12ms/step - loss: 1.0408 - accura
cy: 0.6458
Epoch 8/10
547/547 [=====] - 7s 12ms/step - loss: 0.9479 - accura
cy: 0.6789
Epoch 9/10
547/547 [=====] - 7s 12ms/step - loss: 0.8661 - accura
cy: 0.7111
Epoch 10/10
547/547 [=====] - 7s 12ms/step - loss: 0.7642 - accura
cy: 0.7450
```

```
In [11]: history = pd.DataFrame(history.history)
```


```
In [12]: history
```

```
Out[12]:
```

	loss	accuracy
0	1.762730	0.356857
1	1.527228	0.451486
2	1.406960	0.505371
3	1.314531	0.535429
4	1.227925	0.571257
5	1.140598	0.604514
6	1.057282	0.639543
7	0.956235	0.673429
8	0.872383	0.707143
9	0.776495	0.740057

```
In [13]: history.plot(figsize=(16,10),)
```

```
Out[13]: <matplotlib.axes._subplots.AxesSubplot at 0x7f0bb612d510>
```



```
In [14]: ypred = np.argmax(AlexNet.predict(xtest_1),axis=1)

ytrain = np.argmax(ytrain,axis=1)
ytrain_new_50 = np.argmax(ytrain_new_50,axis=1)
ytrain_1 = np.argmax(ytrain_1,axis=1)

ytest = np.argmax(ytest,axis=1)
ytest_new_50 = np.argmax(ytest_new_50,axis=1)
ytest_1 = np.argmax(ytest_1,axis=1)
```

```
In [15]: from sklearn.metrics import confusion_matrix,classification_report,accuracy_score
accuracy_score(ypred,ytest_1)
```

Out[15]: 0.4696

```
In [16]: ytrain = to_categorical(ytrain,10)
ytrain_new_50 = to_categorical(ytrain_new_50,10)
ytrain_1 = to_categorical(ytrain_1,10)

ytest = to_categorical(ytest,10)
ytest_new_50 = to_categorical(ytest_new_50,10)
ytest_1 = to_categorical(ytest_1,10)
```

```
In [17]: history = AlexNet.fit(xtrain,ytrain,epochs=10,batch_size=32)
```

```
Epoch 1/10
1563/1563 [=====] - 19s 12ms/step - loss: 1.0515 - acc
uracy: 0.6403
Epoch 2/10
1563/1563 [=====] - 19s 12ms/step - loss: 0.9427 - acc
uracy: 0.6777
Epoch 3/10
1563/1563 [=====] - 19s 12ms/step - loss: 0.8495 - acc
uracy: 0.7101
Epoch 4/10
1563/1563 [=====] - 19s 12ms/step - loss: 0.7463 - acc
uracy: 0.7505
Epoch 5/10
1563/1563 [=====] - 19s 12ms/step - loss: 0.6588 - acc
uracy: 0.7788
Epoch 6/10
1563/1563 [=====] - 19s 12ms/step - loss: 0.5772 - acc
uracy: 0.8105
Epoch 7/10
1563/1563 [=====] - 19s 12ms/step - loss: 0.4984 - acc
uracy: 0.8361
Epoch 8/10
1563/1563 [=====] - 19s 12ms/step - loss: 0.4401 - acc
uracy: 0.8574
Epoch 9/10
1563/1563 [=====] - 19s 12ms/step - loss: 0.3852 - acc
uracy: 0.8753
Epoch 10/10
1563/1563 [=====] - 19s 12ms/step - loss: 0.3414 - acc
uracy: 0.8905
```

```
In [18]: history = pd.DataFrame(history.history)
```


```
In [19]: history
```

```
Out[19]:
```

	loss	accuracy
0	1.051510	0.64032
1	0.942745	0.67774
2	0.849459	0.71012
3	0.746287	0.75052
4	0.658829	0.77884
5	0.577167	0.81052
6	0.498385	0.83610
7	0.440125	0.85744
8	0.385194	0.87528
9	0.341448	0.89046

```
In [20]: history.plot.line(figsize=(16,10))
```

```
Out[20]: <matplotlib.axes._subplots.AxesSubplot at 0x7f0baf790510>
```



```
In [21]: ypred = np.argmax(AlexNet.predict(xtest),axis=1)
```

```
ytrain = np.argmax(ytrain,axis=1)
ytrain_new_50 = np.argmax(ytrain_new_50,axis=1)
ytrain_1 = np.argmax(ytrain_1,axis=1)

ytest = np.argmax(ytest,axis=1)
ytest_new_50 = np.argmax(ytest_new_50,axis=1)
ytest_1 = np.argmax(ytest_1,axis=1)
```

```
In [22]: from sklearn.metrics import confusion_matrix,classification_report,accuracy_score
accuracy_score(ypred,ytest)
```

```
Out[22]: 0.6164
```